

Data Processing in Scale

INST447 Data Source and Manipulation

Wei Ai
aiwei@umd.edu

Think about the following questions:

- Midjourney: What is the most frequent word in prompt in the last hour?
- ChatGPT (OpenAI): How many turns on average does the conversations have, do they differ by users' age?

UserID	Timestamp	Prompt
A001		draw a cat licking paws on a sofa
A002		a dog resting in front of a sofa
A001		...
A003		...

ChatID	UserID	# turns
B001	C001	3
B002		6
B003		1
B004		10

UserID	Age (inferred)
C001	25
C002	50

How would you do this using Pandas?

Word Count with Pandas

- Split sentences to words

```
words = df["Prompt"].str.lower().str.split()
```

- Aggregate and count

```
word_count = words.groupby(lambda x:x).count()
```

- Chaining together

```
df.Prompt.str.split().explode().reset_index().groupby("Prompt").count()\
    .sort_values("index", ascending=False).head(10)
```

UserID	Timestamp	Prompt
A001		draw a cat licking paws on sofa
A002		a dog resting in front of a sofa
A001		...
A003		...

Prompt
draw
a
cat
licking
...

a	10
cat	8
draw	2
...	

What if the DataFrame Exceeds Memory?

- Option 1: Ditch Pandas, streaming Python

```
ctr = Counter()
for line in sys.stdin: # or open("prompts.txt")
    ctr.update(line.lower().split())

print(ctr.most_common(20))
```

- Option 2: Ditch Python, use shell commands and external storage
 - This works surprisingly well, actually.

```
cat prompts.txt | tr -s ' ' '\n' | sort | uniq -c | sort -nr | head
```

- Limited by single thread on single machine

Word Count with Pandas

- Split sentences to words

```
words = df[“Prompt”].str.lower().str.split()
```

- Aggregate and count

```
word_count = words.groupby(lambda x:x).count()
```

draw a cat licking paws on sofa

a dog resting in front of a sofa

draw a cat licking paws on sofa
a dog resting in front of a sofa

Word Count with Pandas


- Split sentences to words

```
words = df[“Prompt”].str.lower().str.split()
```

- Aggregate and count

```
word_count = words.groupby(lambda x:x).count()
```

draw a cat licking paws on sofa
a dog resting in front of a sofa




a	10
cat	1
sofa	2
...	

MapReduce - A Data Parallel Abstraction

- Process a large number of records
- “Do something” to each **Map**
- Group intermediate results
- “Aggregate” intermediate results **Reduce**
- Write final results

MapReduce




- Programmer specifies two functions:
 - **map** ($k_1, v_1 \rightarrow \text{List}[(k_2, v_2)]$)
 - **reduce** ($k_2, \text{List}[v_2] \rightarrow \text{List}[(k_3, v_3)]$)
- All values with the same key are sent to the same reducer
- The execution framework handles everything else...

Word Count with MapReduce

- **Map:** for each record → emit (word, 1)
- **Combiner (local reduce):** sum counts per word within a mapper (cuts shuffle size)
- **Shuffle:** route same words to same reducer
- **Reduce:** sum all counts; output (word, total_count)


```
map(record):  
    for word in tokenize(df.prompt):  
        emit(word, 1)  
  
    combine(word, counts):  
        emit(word, sum(counts))  
  
    reduce(word, counts):  
        emit(word, sum(counts))
```

draw a cat licking paws on sofa
a dog resting in front of a sofa



a	10
cat	1
sofa	2
...	

MapReduce – Beyond Map and Reduce



MapReduce – Caveats with Combiner

- ChatGPT (OpenAI): How many turns on average does the conversations have, do they differ by users' age?

```
map(record):  
    emit(uid, n_turns)  
  
map(record):  
    emit(uid, n_turns)  
  
reduce(uid, [n_turns]):  
    S = sum([n_turns])  
    C = count([n_turns])  
    avg = S/C  
    emit(uid, avg)
```

```
map(record):  
    emit(uid, n_turns)  
  
combiner(uid, [n_turns]):  
    S = sum([n_turns])  
    C = count([n_turns])  
    avg = S/C  
    emit(uid, avg)  
  
reduce(uid, [avgs]):  
    S = sum([avgs])  
    C = count([avgs])  
    avg = S/C  
    emit(uid, avg)
```

```
map(record):  
    emit(uid, (n_turns, 1))  
  
combiner/reduce(mid, partials):  
    S = sum([S for S in partials])  
    C = sum([C for C in partials])  
    emit(uid, (S, C))  
  
#final (after reduce):  
forall:  
    emit(uid, S/C)
```

Will this work?


Will this work?

The Execution Framework (“Runtime”) Handles “Everything Else”...

- Scheduling
 - Assigns workers to map and reduce tasks
- “Data distribution”
 - Moves processes to data
- Synchronization
 - Groups intermediate data
- Errors and faults
 - Detects worker failures and restarts


Everything happens on top of a distributed FS (later)

Distributed Group By in MapReduce



Don't move data to workers... move workers to the data!

Key idea: co-locate storage and compute
[Start up worker on nodes that hold the data](#)



We need a distributed file system for managing this
[GFS \(Google File System\) for Google's MapReduce](#)
[HDFS \(Hadoop Distributed File System\) for Hadoop](#)

GFS: Assumptions

Commodity hardware over “exotic” hardware

Scale “out”, not “up”

High component failure rates

Inexpensive commodity components fail all the time

“Modest” number of huge files

Multi-gigabyte files are common, if not encouraged

Files are write-once, mostly appended to

Logs are a common case

Large streaming reads over random access

Design for high sustained throughput over low latency

GFS: Design Decisions

Files stored as chunks

Fixed size (64MB)

Reliability through replication

Each chunk replicated across 3+ chunkservers

Single master to coordinate access and hold metadata

Simple centralized management

No data caching

Little benefit for streaming reads over large datasets


Simplify the API: not POSIX!

Push many issues onto the client (e.g., data layout)

HDFS = GFS clone (same basic ideas)

An aerial photograph of a massive data center complex. The facility consists of numerous interconnected buildings, parking lots, and roads, all situated in a vast, flat landscape. In the background, a vibrant sunset or sunrise染色了整个天空，从深蓝色到明亮的黄色和橙色。The text is overlaid on the upper portion of the image.

The datacenter *is* the computer!
What's the instruction set?



So you like programming in assembly?



Hadoop is great, but it's really waaaaay too low level!
(circa 2007)

Data-Parallel Dataflow Languages

We have a collection of **records**,
want to apply a bunch of operations
to compute some result

What are the dataflow operators?

Spark

Answer to “What’s beyond MapReduce?”

Brief history:


Developed at UC Berkeley AMPLab in 2009

Open-sourced in 2010


Became top-level Apache project in February 2014

Commercial support provided by DataBricks


Spark vs. Hadoop



MapReduce




Map-like Operations




(Not meant to be exhaustive)

Reduce-like Operations




(Not meant to be exhaustive)

Sort Operations



(Not meant to be exhaustive)

Join-like Operations




(Not meant to be exhaustive)

Spark Word Count

```
val textFile = sc.textFile(args.input())
```

RDD[T] ??



textFile

```
.flatMap(line => tokenize(line))  
.map(word => (word, 1))  
.reduceByKey((x, y) => x + y)  
.saveAsTextFile(args.output())
```

flatMap
f: (T) =>
TraversableOnce[U]

RDD[U]

What's an RDD?


Resilient Distributed Dataset (RDD)

= immutable = partitioned

Wait, so how do you actually do anything?
Developers define *transformations* on RDDs
Framework keeps track of lineage

RDD Lifecycle

Transformation




Transformations are lazy:
Framework keeps track of lineage

Actions trigger actual execution


Spark Word Count

RDDs



RDDs and Lineage

On HDFS




*Remember,
transformations are lazy!*

Action!


RDDs and Optimizations

Lazy evaluation creates optimization opportunities




Physical Operators

Narrow Dependencies:




map, filter




union


Wide Dependencies:



join with inputs
co-partitioned




groupByKey



join with inputs not
co-partitioned

Execution Plan



Kinda like a sequence of MapReduce jobs

Can I abstract away RDD? I Like DataFrame

```
wc = (prompts_df.select("Prompt").rdd  
      .flatMap(lambda r: (r["Prompt"] or "").split())  
      .map(lambda w: (w, 1))  
      .reduceByKey(lambda a, b: a + b))
```

```
from pyspark.sql.functions import col, split, explode
```

```
wc = (prompts_df  
      .withColumn("word", explode(split(col("Prompt"), " ")))  
      .groupBy("word").count())
```

```
df.Prompt.str.split().explode().reset_index().groupby("Prompt").count()
```

How to Mapreduce Sandwiches

