

Information Retrieval Evolution

INST447 Data Source and Manipulation

Wei Ai

What is Information Retrieval?

- **Narrow Definition:** Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers).
 - *Manning, Raghavan, and Schütze*
- (We will use Text Retrieval and Information Retrieval interchangeably for this lecture)

Broader Definitions

- Data can be semi-structured, multi-modal, abstract.
 - Image, video, opinion, expert, ...
 - Although in this class, we will only talk about text
- Information need can be implicit, dynamic, inaccurate
 - Information filtering, recommender systems
 - Sometimes there isn't a query!
- Find information → knowledge acquisition
 - Give me what you have → tell me what you know
- Relevance isn't the only criterion.
 - Novelty, diversity, personalization, ...

Short vs. Long Term Info Need

- **Short-term** information need (Ad hoc retrieval)
 - “Temporary need”, e.g., info about used cars
 - Information source is relatively static
 - User “pulls” information
 - Application example: library search, Web search
- **Long-term** information need (Filtering)
 - “Stable need”, e.g., new data mining algorithms
 - Information source is dynamic
 - System “pushes” information to user
 - Applications: news filter, recommender systems

Era 1: Boolean Retrieval

- Documents and queries as **sets** of terms;
- Matching/relevance via set operations: AND/OR/NOT
- Example:
 - (data AND science) OR (machine AND learning) NOT statistics
- Fast, interpretable, controllable; fragile to vocabulary mismatch
- Still everywhere (filters, fielded search, grep, SQL LIKE)

Era 1: Boolean Retrieval

```
(  
    "Vaccination Hesitancy" [Mesh]  
    OR (vaccin*[tiab] AND (hesitan*[tiab] OR accept*[tiab] OR refusal*[tiab] OR uptake[tiab] OR intend*[tiab]))  
)  
AND  
(  
    "Health Personnel" [Mesh]  
    OR healthcare worker*[tiab] OR health care worker*[tiab] OR "medical staff" [tiab]  
    OR nurse*[tiab] OR physician*[tiab] OR clinician*[tiab] OR doctor*[tiab]  
)  
AND  
(  
    "Surveys and Questionnaires" [Mesh]  
    OR Survey [Publication Type]  
    OR survey*[tiab] OR questionnair*[tiab] OR "cross-sectional" [tiab] OR cross sectional[tiab] OR prevalence[tiab]  
)  
AND humans [Mesh]  
NOT animals [Mesh]  
NOT (Editorial [ptyp] OR Letter [ptyp] OR Comment [ptyp] OR News [ptyp])  
AND english [lang]  
AND ("2019/01/01" [dp] : "3000" [dp])
```

Beyond Boolean Retrieval

- Given a query, how do we know if document A is **more** relevant than B?
- One Possible Answer:
If document A uses more query words than document B
(Word usage in document A is more similar to that in query)

Revisit: Vector Representation and Similarity

- Represent a doc/query by a term vector
 - Term: **basic concept**, e.g., word or phrase
 - Each term defines one dimension
 - N terms define a high-dimensional space
 - Element of vector corresponds to term weight
 - E.g., $d = (x_1, \dots, x_N)$, x_i is “**importance**” of term i
- Measure relevance by the similarity/distance (e.g., the cosine similarity) between the query vector and document vector in the vector space

What the VS model doesn't say

- How to define/select the “basic concept”
 - How to select index terms
 - Concepts are assumed to be orthogonal
- How to assign weights
 - Weight in query indicates importance of term
 - Weight in doc indicates how well the term characterizes the document
 - We talked about simple presence/absence
- How to define the similarity/distance measure

Era 2(a): Vector Space (Sparse Vectors)

- Key techniques:
 - Vector Space Model: Documents and queries as vectors
 - Rank documents by vector similarity
 - TF-IDF: Term frequency \times Inverse document frequency
 - How important is a word to *this* doc?
- Relevance = Similarity

Term Frequency (TF) Weighting

- Idea: A term is more important if it occurs more frequently in a document
- Formulas:
 - $c(t,d)$: the frequency count of term t in doc d
 - Raw TF: $TF(t, d) = c(t, d)$
- We always need to normalize the raw TF
 - “Repeated occurrences” are less informative than the “first occurrence”
 - Log TF: $TF(t, d) = \log(c(t, d) + 1)$

Inverted Document Frequency (IDF) Weighting

- Idea: A term is more discriminative if it only occurs in fewer documents.
 - Why is this true?
- Formula:
 - $IDF(t) = 1 + \log\left(\frac{n}{k}\right)$
 - N - total number of documents in collection
 - K - number of document with term t
(document frequency)

Note that IDF is document independent while TF is document dependent!


Era 2(b): Probabilistic Models

- "What's the *probability* this doc is relevant?"
- Relevance = Probability
- BM25 - the “king” of baselines
 - A "bag of words" model, but smarter:
 - it accounts for term saturation (the 10th "data" is less important than the 1st)
 - and document length – penalizing longer document.

$$\text{score}(D, Q) = \sum_{i=1}^n \text{IDF}(q_i) \cdot \frac{f(q_i, D) \cdot (k_1 + 1)}{f(q_i, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgdl}}\right)}$$

Era 3: Graph Signals

- Vector similarity isn't enough
 - "Best Python tutorial" returns random blog posts
 - Need to know which sources are authoritative
- PageRank and variants: use the web graph to estimate authority
- The Pipeline
 - Retrieve: Use vector methods (BM25) to find candidate documents
 - Rank: Combine text relevance + PageRank score
 - Present: Top results balance relevance AND authority



Era 4: Learning to Rank (LTR)

- Why use just one or two signals? Let's use hundreds and “learn” the best ranking formula.
- Assembling multiple features:
 - BM25 scores, PageRank score, user behavior, freshness, ...
- Pointwise/pairwise/listwise training with editorial/click data
- Classical: RankSVM, Gradient Boosted Trees (LambdaMART)
- Often used as re-ranker over a first-pass retrieve.

Era 5: Neural Information Retrieval

- From Sparse to Dense
- The shift: Vectors become embeddings
 - **Sparse Vector:** 50,000+ dimensions, full of zeros. [0, 0, 1, 0, ..., 0, 2, 0]
 - **Dense Vector:** 300-1000+ dimensions, all have meaning.[0.23, -0.45, 1.03, ..., 0.91]
 - Capture semantic similarity: "car" close to "automobile"
- Key technologies:
 - Word2Vec, GloVe (word embeddings)
 - BERT, RoBERTa (contextual embeddings)
 - Sentence transformers (document embeddings)
- Retrieval function is now a deep neural network). You find documents that are semantically close in this embedding space, even if not sharing keywords.

Vector Database

- Traditional Databases are built for exact matches and filtering on structured data.
 - e.g. `SELECT * FROM docs WHERE author = 'Smith'`
 - Terrible at finding "nearest neighbors" in 1000-dimensional space. A `WHERE` clause can't do that.
- Problem: $10M \text{ documents} \times 768 \text{ dimensions} = 7.68 \text{ billion numbers}$
 - Can't do linear search (too slow!)
 - Traditional SQL databases are terrible at finding "nearest neighbors" in such dense vector space.
- Solution: Vector Databases (Pinecone, Chroma, FAISS)
 - Use approximate nearest neighbor search (ANN)
 - Trade perfect accuracy for speed

RAG – Retrieval-Augmented Generation

- The pipeline:
 - Retrieve: Find relevant documents (using methods we discussed)
 - Query understanding: expand/rewrite, detect intent
 - Retrieve: BM25/sparse, dense, or hybrid
 - Re-rank: cross-encoder for top-k quality
 - Augment: Add retrieved docs to LLM prompt
 - Generate: LLM synthesizes answer with citations
 - Post-process: guardrails, citation verification, caching
- Why RAG?
 - Ground LLMs with external knowledge,
 - reduce hallucination, add freshness.
 - Allows dynamic knowledge updates

The Agentic AI Twist

```
grep -r "TODO" ./src  
find . -name "*.py" -exec grep "def process_data" {} \;
```

- LLM agents often skip RAG, or even dense vectors.
- Observation: Claude Code, Devin, Cursor use grep and shell commands
- Why?
 - Speed: grep is instant, no GPU needed
 - Structured data: Code has structure (function names, file paths)
 - Exact matching: When you want "error code 404", not "similar to error"
 - Transparency: Can explain what was searched
- Different data genres need different retrieval methods!
 - Unstructured text → embeddings (semantic search)
 - Structured code → regex/grep (exact match)
 - Logs → SQL (structured queries)
 - Knowledge graphs → graph traversal
- Know your data type, choose your representation!

Long-Context LLMs: IR Killer?

- The dream: 1M+ token context
 - Dump all documents into context
 - Ask questions
 - No retrieval needed!
- The reality check:
 - "Needle in haystack" problem: LLMs still miss info in long contexts
 - Cost: Processing 1M tokens is expensive
 - Latency: Takes time to process
 - The core IR problem remains: Users still need help finding what matters

Long Context ≠ No Retrieval

- Even with infinite context...
- You still need IR for:
 - Filtering: What goes into the context?
 - Prioritization: What appears first? (LLMs attend to start/end more)
 - Summarization: Can't show user 1M tokens
 - Cost optimization: Only retrieve what's needed
- The constant: Understanding and satisfying information needs without overwhelming users.

Evaluation: What to Measure (and Why)

- Search is not classification
- Scenario: Search for "machine learning tutorials"
 - 1 million relevant documents exist
 - System returns 10 documents
- Questions:
 - Did we find good ones? → Precision
 - Did we find all of them? → Recall
 - Are the best ones at the top? → NDCG, MAP
- Other practical metrics:
 - Diversity/novelty, coverage; latency; cost


Evaluation in the LLM Era

- New metrics for new tasks.
- Classical IR: Rank documents
- LLM + RAG: Generate answers
 - Faithfulness: Did LLM stick to retrieved docs?
 - Attribution: Can we trace claims to sources?
 - Answer quality: Is it correct, complete, concise?
- The evolution: Task changed, metrics evolved, but still measuring "information need satisfaction"

Looking Ahead: IR to serve humans and AI agents

- IR to serve human users: from search engines to task agents
 - Human-centered design
 - Information seeking behavior / intent modeling / preference learning...
- IR to serve AI agents: five essential new IR tasks
 - LLM for IR agents (optimize the pipeline)
 - Technical: indexing, retrieval models, re-ranking, vector DB ops, evaluation, ...

Intelligence of Information Retrieval Systems



Safety, Privacy, and Integrity

- Prompt injection & data exfiltration in RAG/agents
- Copyright & licensing of indexed sources
- PII/PHI handling; on-device retrieval, encryption, access control
- Guardrails: allowlists, input/output filters, sandboxed tools,...

When Working on IR Problems

- Ask these questions:
 - What is the data type? (text, code, logs, structured)
 - What is the scale? (1K docs vs 1B docs)
 - What is the need? (exact match vs semantic search)
 - What are the constraints? (latency, cost, accuracy)
- Then choose:
 - Exact string matching (grep, SQL)
 - Sparse retrieval (BM25)
 - Dense retrieval (embeddings)
 - Graph-based (PageRank)
 - Hybrid (usually the answer!)

Case Study: Build a Health/Medical Search Engine

Question	Answer	Implication
Data type?	Scientific text + structured drug data	Need both semantic understanding AND exact matching
Scale?	34M+ documents	Need fast first-pass filtering
User need?	Doctors: precise terms / Patients: natural language	Must serve different query styles
Constraints?	Accuracy critical, must explain results	False info = harm; need source attribution

Case Study: Build a Corporate Knowledge Base

Question	Answer	Implication
Data type?	Text docs + code + messages + structured tickets	Different sources need different methods
Scale?	100K-1M items, 1000+ updates/day	Moderate scale, high freshness need
User need?	Known-item ("Q3 report") + exploratory ("our AI work")	Mix of exact and semantic
Constraints?	Privacy critical , cost-conscious, <500ms latency	Permissions first! Budget matters

Takeaway

- Information Retrieval is about understanding and meeting information needs.
- The representations change
 - sets → sparse vectors → dense vectors → ?
 - graph, stream, ...
- The core mission stays
 - find the right information, don't overwhelm
- Your role as data professionals:
 - Choose the right tool for the right job,
 - understand the tradeoffs,
 - and keep learning!

If You Want to Build Your Own IR System

- Classic IR:
 - Elasticsearch: Industry standard
 - Apache Solr: Open source alternative
- RAG Frameworks:
 - LangChain: Full-featured, lots of integrations
 - LlamaIndex: Focused on RAG
- Vector Databases (free tiers available):
 - Chroma: Local, easy to start
 - Pinecone: Cloud, generous free tier
- Evaluation:
 - TREC datasets
- Start small, iterate, learn by building!